

深入理解Spark 2.1 Core （四）：运算结果处理和容错的原理与源码分析

版权声明：本文为博主原创文章，转载请附上原文地址。 <http://blog.csdn.net/u011239443/article/details/54024946>

在上一篇博文《深入理解Spark 2.1 Core （三）：任务调度器的实现与源码分析》TaskScheduler在发送任务给executor前的工作就全部完成了。这篇博文，我们来看看当executor计算完任务后，Spark是如何处理获取的计算结果与容错的。

调用栈如下：

- TaskSchedulerImpl.statusUpdate
 - TaskResultGetter.enqueueSuccessfulTask
 - TaskSchedulerImpl.handleSuccessfulTask
 - TaskSetManager.handleSuccessfulTask
 - DAGScheduler.taskEnded
 - DAGSchedulerEventProcessLoop.doOnReceive
 - DAGScheduler.handleTaskCompletion
 - TaskResultGetter.enqueueFailedTask
 - TaskSchedulerImpl.handleFailedTask
 - TaskSetManager.handleFailedTask
 - DAGScheduler.taskEnded
 - DAGSchedulerEventProcessLoop.doOnReceive
 - DAGScheduler.handleTaskCompletion

TaskSchedulerImpl.statusUpdate

TaskRunner将任务的执行结果发送给DriverEndPoint，DriverEndPoint会转给TaskSchedulerImpl的statusUpdate：

```
1 def statusUpdate(tid: Long, state: TaskState, serializedData: ByteBuffer) {
2   var failedExecutor: Option[String] = None
3   var reason: Option[ExecutorLossReason] = None
4   synchronized {
5     try {
6       taskIdToTaskSetManager.get(tid) match {
7         case Some(taskSet) =>
8           //这只针对Mesos调度模式
9           if (state == TaskState.LOST) {
10             val execId = taskIdToExecutorId.getOrElse(tid, throw new IllegalStateException(
11               "taskIdToTaskSetManager.contains(tid) <=> taskIdToExecutorId.contains(tid)"))
12             if (executorIdToRunningTaskIds.contains(execId)) {
13               reason = Some(
14                 SlaveLost(s"Task $tid was lost, so marking the executor as lost as well."))
15               removeExecutor(execId, reason.get)
16               failedExecutor = Some(execId)
17             }
18           }
19           //FINISHED KILLED LOST 都属于 isFinished
20           if (TaskState.isFinished(state)) {
21             cleanupTaskState(tid)
```

```

22      taskSet.removeRunningTask(tid)
23      //若FINISHED调用taskResultGetter.enqueueSuccessfulTask,
24      //否则调用taskResultGetter.enqueueFailedTask(taskSet, tid, state, serializedData)
25      if (state == TaskState.FINISHED) {
26          taskResultGetter.enqueueSuccessfulTask(taskSet, tid, serializedData)
27      } else if (Set(TaskState.FAILED, TaskState.KILLED, TaskState.LOST).contains(state)) {
28          taskResultGetter.enqueueFailedTask(taskSet, tid, state, serializedData)
29      }
30  }
31  case None =>
32      logError(
33          ("Ignoring update with state %s for TID %s because its task set is gone (this is " +
34           "likely the result of receiving duplicate task finished status updates) or its " +
35           "executor has been marked as failed.")
36              .format(state, tid))
37      }
38  } catch {
39      case e: Exception => logError("Exception in statusUpdate", e)
40  }
41  }
42  if (failedExecutor.isDefined) {
43      assert(reason.isDefined)
44      dagScheduler.executorLost(failedExecutor.get, reason.get)
45      backend.reviveOffers()
46  }
47  }

```

处理执行成功的结果

我们先看下处理执行成功的结果的运行机制：

TaskResultGetter.enqueueSuccessfulTask

```

1  def enqueueSuccessfulTask(
2      taskSetManager: TaskSetManager,
3      tid: Long,
4      serializedData: ByteBuffer): Unit = {
5      //通过线程池来获取结果
6      getTaskResultExecutor.execute(new Runnable {
7          override def run(): Unit = Utils.logUncaughtExceptions {
8              try {
9                  val (result, size) = serializer.get().deserialize[TaskResult[_]](serializedData) match {
10                     //可以直接获取到的结果
11                     case directResult: DirectTaskResult[_] =>
12                         //判断大小是否符合要求
13                         if (!taskSetManager.canFetchMoreResults(serializedData.limit())) {
14                             return
15                         }
16
17                     directResult.value(taskResultSerializer.get())
18                     (directResult, serializedData.limit())
19                     //若不能直接获取到结果
20                     case IndirectTaskResult(blockId, size) =>
21                         if (!taskSetManager.canFetchMoreResults(size)) {
22                             // 判断大小是否符合要求,
23                             //若不符合则远程的删除计算结果
24                             sparkEnv.blockManager.master.removeBlock(blockId)
25                             return
26                         }
27                         logDebug("Fetching indirect task result for TID %s".format(tid))
28                         scheduler.handleTaskGettingResult(taskSetManager, tid)
29                         val serializedTaskResult = sparkEnv.blockManager.getRemoteBytes(blockId)
30                         //从远程获取计算结果

```

```

31         if (!serializedTaskResult.isDefined) {
32             //若在任务执行结束后与我们去获取结果之间机器出现故障了
33             //或者block manager 不得不刷新结果了
34             //那么我们将不能够获取到结果
35             scheduler.handleFailedTask(
36                 taskSetManager, tid, TaskState.FINISHED, TaskResultLost)
37             return
38         }
39         val deserializedResult = serializer.get().deserialize[DirectTaskResult[_]](
40             serializedTaskResult.get.toByteBuffer)
41         // 反序列化
42         deserializedResult.value(taskResultSerializer.get())
43         sparkEnv.blockManager.master.removeBlock(blockId)
44         (deserializedResult, size)
45     }
46
47
48     result.accumUpdates = result.accumUpdates.map { a =>
49         if (a.name == Some(InternalAccumulator.RESULT_SIZE)) {
50             val acc = a.asInstanceOf[LongAccumulator]
51             assert(acc.sum == 0L, "task result size should not have been set on the executors")
52             acc.setValue(size.toLong)
53             acc
54         } else {
55             a
56         }
57     }
58
59     //处理获取到的计算结果
60     scheduler.handleSuccessfulTask(taskSetManager, tid, result)
61 } catch {
62     case cnf: ClassNotFoundException =>
63         val loader = Thread.currentThread.getContextClassLoader
64         taskSetManager.abort("ClassNotFoundException with classloader: " + loader)
65     case NonFatal(ex) =>
66         logError("Exception while getting task result", ex)
67         taskSetManager.abort("Exception while getting task result: %s".format(ex))
68 }
69 }
70 })
71 }

```

TaskSchedulerImpl.handleSuccessfulTask

调用taskSetManager.handleSuccessfulTask

```

1  def handleSuccessfulTask(
2      taskSetManager: TaskSetManager,
3      tid: Long,
4      taskResult: DirectTaskResult[_]): Unit = synchronized {
5      taskSetManager.handleSuccessfulTask(tid, taskResult)
6  }

```

TaskSetManager.handleSuccessfulTask

```

1  def handleSuccessfulTask(tid: Long, result: DirectTaskResult[_]): Unit = {
2      val info = taskInfos(tid)
3      val index = info.index
4      info.markFinished(TaskState.FINISHED)
5      //从RunningTask中移除该task
6      removeRunningTask(tid)
7      //通知dagScheduler该task完成
8      sched.dagScheduler.taskEnded(tasks(index), Success, result.value(), result.accumUpdates, info)

```

```

9      //杀死所有其他与之相同的task的尝试
10     for (attemptInfo <- taskAttempts(index) if attemptInfo.running) {
11         logInfo(s"Killing attempt ${attemptInfo.attemptNumber} for task ${attemptInfo.id} " +
12             s"in stage ${taskSet.id} (TID ${attemptInfo.taskId}) on ${attemptInfo.host} " +
13             s"as the attempt ${info.attemptNumber} succeeded on ${info.host}")
14         sched.backend.killTask(attemptInfo.taskId, attemptInfo.executorId, true)
15     }
16     if (!successful(index)) {
17         //计数
18         tasksSuccessful += 1
19         logInfo(s"Finished task ${info.id} in stage ${taskSet.id} (TID ${info.taskId}) in" +
20             s" ${info.duration} ms on ${info.host} (executor ${info.executorId})" +
21             s" ($tasksSuccessful/$numTasks)")
22         //若果有所task成功了,
23         //那么标记successful, 并且停止
24         successful(index) = true
25         if (tasksSuccessful == numTasks) {
26             isZombie = true
27         }
28     } else {
29         logInfo("Ignoring task-finished event for " + info.id + " in stage " + taskSet.id +
30             " because task " + index + " has already completed successfully")
31     }
32     maybeFinishTaskSet()
33 }

```

DAGScheduler.taskEnded

我们再深入看下是如何通知dagScheduler该task完成的：

```

1  def taskEnded(
2      task: Task[_],
3      reason: TaskEndReason,
4      result: Any,
5      accumUpdates: Seq[AccumulatorV2[_], _]],
6      taskInfo: TaskInfo): Unit = {
7      //发送CompletionEvent信号
8      eventProcessLoop.post(
9          CompletionEvent(task, reason, result, accumUpdates, taskInfo))
10 }

```

DAGSchedulerEventProcessLoop.doOnReceive

上一篇博文讲过，DAGSchedulerEventProcessLoop的doOnReceive会对信号进行监听：

```

1      case completion: CompletionEvent =>
2          dagScheduler.handleTaskCompletion(completion)

```

DAGScheduler.handleTaskCompletion

我们来看下DAGScheduler.handleTaskCompletion部分核心代码：

```

1  ***
2      //根据stageId 得到stage
3      val stage = stageIdToStage(task.stageId)
4      //这里的event就是completion
5      event.reason match {
6          //这里只看成功的流程
7          case Success =>
8              //将这个task 从stage等待处理分区中删去
9              stage.pendingPartitions -= task.partitionId
10             task match {

```

```

11 //若是最后一个Stage的task
12 case rt: ResultTask[_] =>
13 //将stage 转为 ResultStage
14 val resultStage = stage.asInstanceOf[ResultStage]
15 resultStage.activeJob match {
16 //获取这Stage的job
17 case Some(job) =>
18     if (!job.finished(rt.outputId)) {
19         updateAccumulators(event)
20         //标记状态
21         job.finished(rt.outputId) = true
22         //计数
23         job.numFinished += 1
24         // 若Job的所有partition都完成了,
25         //移除这个Job
26         if (job.numFinished == job.numPartitions) {
27             markStageAsFinished(resultStage)
28             cleanupStateForJobAndIndependentStages(job)
29             listenerBus.post(
30                 SparkListenerJobEnd(job.jobId, clock.getTimeMillis(), JobSucceeded))
31         }
32         //通知 JobWaiter 有任务成功
33         //但 taskSucceeded 会运行用户自定义的代码
34         //因此可能抛出异常
35         try {
36             job.listener.taskSucceeded(rt.outputId, event.result)
37         } catch {
38             case e: Exception =>
39                 // 标记为失败
40                 job.listener.jobFailed(new SparkDriverExecutionException(e))
41         }
42     }
43     case None =>
44         logInfo("Ignoring result from " + rt + " because its job has finished")
45 }
46 //若不是最后一个Stage的Task
47 case smt: ShuffleMapTask =>
48     val shuffleStage = stage.asInstanceOf[ShuffleMapStage]
49     updateAccumulators(event)
50     val status = event.result.asInstanceOf[MapStatus]
51     val execId = status.location.executorId
52     logDebug("ShuffleMapTask finished on " + execId)
53     if (failedEpoch.contains(execId) && smt.epoch <= failedEpoch(execId)) {
54         logInfo(s"Ignoring possibly bogus $smt completion from executor $execId")
55     } else {
56         //将Task的partitionId和status
57         //追加到OutputLoc
58         shuffleStage.addOutputLoc(smt.partitionId, status)
59     }
60
61     if (runningStages.contains(shuffleStage) && shuffleStage.pendingPartitions.isEmpty) {
62         markStageAsFinished(shuffleStage)
63         logInfo("looking for newly runnable stages")
64         logInfo("running: " + runningStages)
65         logInfo("waiting: " + waitingStages)
66         logInfo("failed: " + failedStages)
67
68         //将outputLoc信息注册到mapOutputTracker
69         //上篇博文中有提到:
70         //首先ShuffleMapTask的计算结果(其实是计算结果数据所在的位置、大小等元数据信息)都会传给Driver的mapOutputTracker
71         // 所以 DAGScheduler.newOrUsedShuffleStage需要先判断Stage是否已经被计算过
72         ///若计算过, DAGScheduler.newOrUsedShuffleStage则把结果复制到新创建的stage
73         //如果没计算过, DAGScheduler.newOrUsedShuffleStage就向注册mapOutputTracker Stage, 为存储元数据占位
74         mapOutputTracker.registerMapOutputs(

```

```

75         shuffleStage.shuffleDep.shuffleId,
76         shuffleStage.outputLocInMapOutputTrackerFormat(),
77         changeEpoch = true)
78
79     clearCacheLocs()
80
81     if (!shuffleStage.isAvailable) {
82         //若Stage不可用（一些任务失败），则从新提交Stage
83         logInfo("Resubmitting " + shuffleStage + " (" + shuffleStage.name +
84             ") because some of its tasks had failed: " +
85             shuffleStage.findMissingPartitions().mkString(", "))
86         submitStage(shuffleStage)
87     } else {
88         // 若该Stage的所有分区都完成了
89         if (shuffleStage.mapStageJobs.nonEmpty) {
90             val stats = mapOutputTracker.getStatistics(shuffleStage.shuffleDep)
91             //将各个Task的标记为Finished
92             for (job <- shuffleStage.mapStageJobs) {
93                 markMapStageJobAsFinished(job, stats)
94             }
95         }
96         //提交该Stage的正在等待的Child Stages
97         submitWaitingChildStages(shuffleStage)
98     }
99 }
100 }
101 ***

```

处理执行失败的结果

TaskResultGetter.enqueueFailedTask

下面，我们回归头来看如何处理失败的结果。

```

1  def enqueueFailedTask(taskSetManager: TaskSetManager, tid: Long, taskState: TaskState,
2      serializedData: ByteBuffer) {
3      var reason : TaskFailedReason = UnknownReason
4      try {
5          //通过线程池来处理结果
6          getTaskResultExecutor.execute(new Runnable {
7              override def run(): Unit = Utils.logUncaughtExceptions {
8                  val loader = Utils.getContextOrSparkClassLoader
9                  try {
10                     //若序列化数据，即TaskFailedReason，存在且长度大于0
11                     //则反序列化获取它
12                     if (serializedData != null && serializedData.limit() > 0) {
13                         reason = serializer.get().deserialize[TaskFailedReason](
14                             serializedData, loader)
15                     }
16                 } catch {
17                     //若是ClassNotFoundException,
18                     //打印log
19                     case cnd: ClassNotFoundException =>
20                         logError(
21                             "Could not deserialize TaskEndReason: ClassNotFoundException with classloader " + loader)
22                     //若其他异常,
23                     //不进行操作
24                     case ex: Exception =>
25                 }
26                 //处理失败的任务
27                 scheduler.handleFailedTask(taskSetManager, tid, taskState, reason)
28             }
29         })

```

```

30     } catch {
31         case e: RejectedExecutionException if sparkEnv.isStopped =>
32     }
33 }

```

TaskSchedulerImpl.handleFailedTask

```

1  def handleFailedTask(
2      taskSetManager: TaskSetManager,
3      tid: Long,
4      taskState: TaskState,
5      reason: TaskFailedReason): Unit = synchronized {
6      //处理失败任务
7      taskSetManager.handleFailedTask(tid, taskState, reason)
8      if (!taskSetManager.isZombie && taskState != TaskState.KILLED) {
9          //handleFailedTask会将失败任务放入待运行的队列等待下一次调度
10         //所以这里开始新一轮调度
11         backend.reviveOffers()
12     }
13 }

```

TaskSetManager.handleFailedTask

我们来看下handleFailedTask核心代码：

```

1  ***
2      //调用dagScheduler处理失败任务
3      sched.dagScheduler.taskEnded(tasks(index), reason, null, accumUpdates, info)
4
5      if (successful(index)) {
6          logInfo(
7              s"Task ${info.id} in stage ${taskSet.id} (TID $tid) failed, " +
8              "but another instance of the task has already succeeded, " +
9              "so not re-queuing the task to be re-executed.")
10     } else {
11         //将这个任务重新加入到等待队列中
12         addPendingTask(index)
13     }
14
15     if (!isZombie && reason.countTowardsTaskFailures) {
16         taskSetBlacklistHelperOpt.foreach(_.updateBlacklistForFailedTask(
17             info.host, info.executorId, index))
18         assert (null != failureReason)
19         //计数 这个任务的重试次数
20         numFailures(index) += 1
21         //若大于等于最大重试次数，默认为4，
22         //则取消这个任务
23         if (numFailures(index) >= maxTaskFailures) {
24             logError("Task %d in stage %s failed %d times; aborting job".format(
25                 index, taskSet.id, maxTaskFailures))
26             abort("Task %d in stage %s failed %d times, most recent failure: %s\nDriver stacktrace:"
27                 .format(index, taskSet.id, maxTaskFailures, failureReason), failureException)
28             return
29         }
30     }
31     maybeFinishTaskSet()
32 }

```

DAGScheduler.handleTaskCompletion

与处理成功结果的过程相同，接下来也会调用DAGScheduler.taskEnded。DAGSchedulerEventProcessLoop的doOnReceive接收CompletionEvent信号，调用dagScheduler.handleTaskCompletion(completion)

我们来看下DAGScheduler.handleTaskCompletion 处理失败任务部分的核心代码：

```

1      //重新提交任务
2      case Resubmitted =>
3          logInfo("Resubmitted " + task + ", so marking it as still running")
4          //把任务加入的等待队列
5          stage.pendingPartitions += task.partitionId
6
7      //获取结果失败
8      case FetchFailed(bmAddress, shuffleId, mapId, reduceId, failureMessage) =>
9          val failedStage = stageIdToStage(task.stageId)
10         val mapStage = shuffleIdToMapStage(shuffleId)
11         //若失败的尝试ID 不是 stage尝试ID,
12         //则忽略这个失败
13         if (failedStage.latestInfo.attemptId != task.stageAttemptId) {
14             logInfo(s"Ignoring fetch failure from $task as it's from $failedStage attempt" +
15                 s" ${task.stageAttemptId} and there is a more recent attempt for that stage " +
16                 s"(attempt ID ${failedStage.latestInfo.attemptId}) running")
17         } else {
18             //若失败的Stage还在运行队列,
19             //标记这个Stage完成
20             if (runningStages.contains(failedStage)) {
21                 logInfo(s"Marking $failedStage (${failedStage.name}) as failed " +
22                     s"due to a fetch failure from $mapStage (${mapStage.name})")
23                 markStageAsFinished(failedStage, Some(failureMessage))
24             } else {
25                 logDebug(s"Received fetch failure from $task, but its from $failedStage which is no " +
26                     s"longer running")
27             }
28             //若不允许重试,
29             //则停止这个Stage
30             if (disallowStageRetryForTest) {
31                 abortStage(failedStage, "Fetch failure will not retry stage due to testing config",
32                     None)
33             }
34             //若达到最大重试次数,
35             //则停止这个Stage
36             else if (failedStage.failedOnFetchAndShouldAbort(task.stageAttemptId)) {
37                 abortStage(failedStage, s"$failedStage (${failedStage.name}) " +
38                     s"has failed the maximum allowable number of " +
39                     s"times: ${Stage.MAX_CONSECUTIVE_FETCH_FAILURES}. " +
40                     s"Most recent failure reason: ${failureMessage}", None)
41             } else {
42                 if (failedStages.isEmpty) {
43                     //若失败的Stage中, 没有个task完成了,
44                     //则重新提交Stage。
45                     //若果有完成的task的话, 我们不能重新提交Stage,
46                     //因为有些task已经被调度过了。
47                     //task级别的重新提交是在TaskSetManager.handleFailedTask进行的
48                     logInfo(s"Resubmitting $mapStage (${mapStage.name}) and " +
49                         s"$failedStage (${failedStage.name}) due to fetch failure")
50                     messageScheduler.schedule(new Runnable {
51                         override def run(): Unit = eventProcessLoop.post(ResubmitFailedStages)
52                     }, DAGScheduler.RESUBMIT_TIMEOUT, TimeUnit.MILLISECONDS)
53                 }
54                 failedStages += failedStage
55                 failedStages += mapStage
56             }
57             // 移除OutputLoc中的数据
58             // 取消注册mapOutputTracker
59             if (mapId != -1) {
60                 mapStage.removeOutputLoc(mapId, bmAddress)
61                 mapOutputTracker.unregisterMapOutput(shuffleId, mapId, bmAddress)
62             }

```



```
63
64     //当有executor上发生多次获取结果失败,
65     //则标记这个executor丢失
66     if (bmAddress != null) {
67         handleExecutorLost(bmAddress.executorId, filesLost = true, Some(task.epoch))
68     }
69 }
70
71 //拒绝处理
72 case commitDenied: TaskCommitDenied =>
73     // 不做什么事,
74     //让 TaskScheduler 来决定如何处理
75
76 //异常
77 case exceptionFailure: ExceptionFailure =>
78     // 更新accumulator
79     updateAccumulators(event)
80
81 //task结果丢失
82 case TaskResultLost =>
83     // 不做什么事,
84     // 让 TaskScheduler 处理这些错误和重新提交任务
85
86 // executor 丢失
87 // 任务被杀死
88 // 未知错误
89 case _: ExecutorLostFailure | TaskKilled | UnknownReason =>
90     // 不做什么事,
91     // 若这task不断的错误,
92     // TaskScheduler 会停止 job
```