

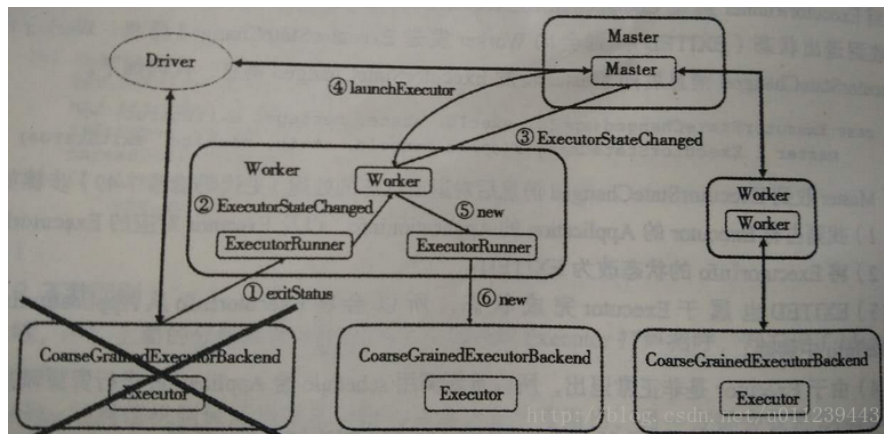
深入理解Spark 2.1 Core（八）：Standalone模式容错及HA的原理与源码分析

版权声明：本文为博主原创文章，转载请附上原文地址。 <http://blog.csdn.net/u011239443/article/details/54287787>

第五、第六、第七篇博文，我们讲解了Standalone模式集群是如何启动的，一个App起来了后，集群是如何分配资源，Worker启动Executor的，Task来是如何执行它，执行得到的结果如何处理，以及app退出后，分配了的资源如何回收。

但在分布式系统中，由于机器众多，所有发生故障是在所难免的，若运行过程中Executor、Worker或者Master异常退出了，那该怎么办呢？这篇博文，我们就来讲讲在Standalone模式下，Spark的集群容错与高可用性（HA）。

Executor



Worker.receive

我先回到《深入理解Spark 2.1 Core（六）：资源调度的原理与源码分析》的ExecutorRunner.fetchAndRunExecutor中，看看executor的退出：

```
1 // executor会退出并返回0或者非0的exitCode
2 val exitCode = process.waitFor()
3 state = ExecutorState.EXITED
4 val message = "Command exited with code " + exitCode
5 // 给Worker发送ExecutorStateChanged信号
6 worker.send(ExecutorStateChanged(appId, execId, state, Some(message), Some(exitCode)))
```

worker接收到了ExecutorStateChanged信号后，调用handleExecutorStateChanged

```
1 case executorStateChanged @ ExecutorStateChanged(appId, execId, state, message, exitStatus) =>
2   handleExecutorStateChanged(executorStateChanged)
```

Worker.handleExecutorStateChanged

```
1 private[worker] def handleExecutorStateChanged(executorStateChanged: ExecutorStateChanged):
2   Unit = {
3     // 给Master发送executorStateChanged信号
4     sendToMaster(executorStateChanged)
5     val state = executorStateChanged.state
6     if (ExecutorState.isFinished(state)) {
7       // 释放executor资源
8       val appId = executorStateChanged.appId
9       val fullId = appId + "/" + executorStateChanged.execId
10      val message = executorStateChanged.message
11      val exitStatus = executorStateChanged.exitStatus
```

```

12  executors.get(fullId) match {
13      case Some(executor) =>
14          logInfo("Executor " + fullId + " finished with state " + state +
15              message.map(" message " + _).getOrElse("") +
16              exitStatus.map(" exitStatus " + _).getOrElse(""))
17          executors -= fullId
18          finishedExecutors(fullId) = executor
19          trimFinishedExecutorsIfNecessary()
20          coresUsed -= executor.cores
21          memoryUsed -= executor.memory
22      case None =>
23          logInfo("Unknown Executor " + fullId + " finished with state " + state +
24              message.map(" message " + _).getOrElse("") +
25              exitStatus.map(" exitStatus " + _).getOrElse(""))
26  }
27  maybeCleanupApplication(appId)
28  }
29  }
30  }

```

Master.receive

Master接收到ExecutorStateChanged信号后:

```

1  case ExecutorStateChanged(appId, execId, state, message, exitStatus) =>
2  // 通过appId取到App的信息,
3  // 在App的信息中找到该executor的信息
4  val execOption = idToApp.get(appId).flatMap(app => app.executors.get(execId))
5  execOption match {
6      case Some(exec) =>
7          val appInfo = idToApp(appId)
8          // 改变改executor的状态
9          val oldState = exec.state
10         exec.state = state
11
12         if (state == ExecutorState.RUNNING) {
13             assert(oldState == ExecutorState.LAUNCHING,
14                 s"executor $execId state transfer from $oldState to RUNNING is illegal")
15             appInfo.resetRetryCount()
16         }
17
18         exec.application.driver.send(ExecutorUpdated(execId, state, message, exitStatus, false))
19
20         if (ExecutorState.isFinished(state)) {
21             logInfo(s"Removing executor ${exec.fullId} because it is $state")
22             // 若该app已经结束,
23             // 保持原来的executor信息,
24             // 用于呈现在Web UI上,
25             // 若该app还没结束,
26             // 则从app信息中移除该executor
27             if (!appInfo.isFinished) {
28                 appInfo.removeExecutor(exec)
29             }
30             // 把executor从worker中移除
31             exec.worker.removeExecutor(exec)
32
33             // 获取executor退出状态
34             val normalExit = exitStatus == Some(0)
35             // 若executor退出状态非正常,
36             // 且app重新尝试调度次数到达最大重试次数,
37             // 则移除这个app
38             if (!normalExit
39                 && appInfo.incrementRetryCount() >= MAX_EXECUTOR_RETRIES
40                 && MAX_EXECUTOR_RETRIES >= 0) { // < 0 disables this application-killing path

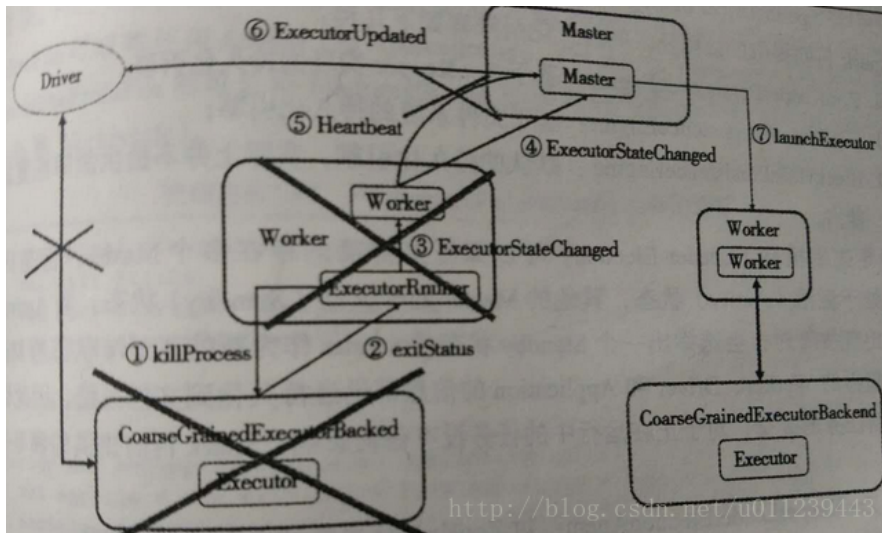
```

```

41     val execs = appInfo.executors.values
42     if (!execs.exists(_.state == ExecutorState.RUNNING)) {
43         logError(s"Application ${appInfo.desc.name} with ID ${appInfo.id} failed " +
44             s"${appInfo.retryCount} times; removing it")
45         removeApplication(appInfo, ApplicationState.FAILED)
46     }
47 }
48 }
49 //重新调度
50 schedule()

```

Worker



Worker.killProcess

我们回到《深入理解Spark 2.1 Core (六) : 资源调度的原理与源码分析》的ExecutorRunner.start中:

```

1 // 创建Shutdownhook线程
2 // 用于worker关闭时, 杀掉executor
3 shutdownHook = ShutdownHookManager.addShutdownHook { () =>
4     if (state == ExecutorState.RUNNING) {
5         state = ExecutorState.FAILED
6     }
7     killProcess(Some("Worker shutting down")) }
8 }

```

worker退出后, ShutdownHookManager会调用killProcess杀死executor:

```

1 private def killProcess(message: Option[String]) {
2     var exitCode: Option[Int] = None
3     if (process != null) {
4         logInfo("Killing process!")
5         // 停止运行日志输出
6         if (stdoutAppender != null) {
7             stdoutAppender.stop()
8         }
9         if (stderrAppender != null) {
10            // 停止错误日志输出
11            stderrAppender.stop()
12        }
13        // kill executor进程,
14        // 并返回结束类型
15        exitCode = Utils.terminateProcess(process, EXECUTOR_TERMINATE_TIMEOUT_MS)

```

```

16     if (exitCode.isEmpty) {
17         logWarning("Failed to terminate process: " + process +
18             ". This process will likely be orphaned.")
19     }
20 }
21 try {
22     // 给worker发送ExecutorStateChanged信号
23     worker.send(ExecutorStateChanged(appId, execId, state, message, exitCode))
24 } catch {
25     case e: IllegalStateException => logWarning(e.getMessage(), e)
26 }
27 }

```

Master.timeOutDeadWorkers

当Worker向Master注册直接的时候，会向worker的handleRegisterResponse发送RegisteredWorker信号。handleRegisterResponse处理该信号时会启动一个线程，来不断的给worker自己的SendHeartbeat信号

```

1     case RegisteredWorker(masterRef, masterWebUiUrl) =>
2     ***
3         forwardMessageScheduler.scheduleAtFixedRate(new Runnable {
4             override def run(): Unit = Utils.tryLogNonFatalError {
5                 self.send(SendHeartbeat)
6             }
7         }, 0, HEARTBEAT_MILLIS, TimeUnit.MILLISECONDS)
8     ***

```

worker receive到SendHeartbeat信号后，处理：

```

1     case SendHeartbeat =>
2         if (connected) { sendToMaster(Heartbeat(workerId, self)) }

```

给Master发送Heartbeat信号。

Master接收到Heartbeat信号后处理：

```

1     case Heartbeat(workerId, worker) =>
2         idToWorker.get(workerId) match {
3             case Some(workerInfo) =>
4                 // 更新worker的lastHeartbeat信息
5                 workerInfo.lastHeartbeat = System.currentTimeMillis()
6             case None =>
7                 if (workers.map(_._id).contains(workerId)) {
8                     logWarning(s"Got heartbeat from unregistered worker $workerId." +
9                         " Asking it to re-register.")
10                    worker.send(ReconnectWorker(masterUrl))
11                } else {
12                    logWarning(s"Got heartbeat from unregistered worker $workerId." +
13                        " This worker was never registered, so ignoring the heartbeat.")
14                }
15        }

```

而在Master.onStart中我可以看到：

```

1     checkForWorkerTimeOutTask = forwardMessageThread.scheduleAtFixedRate(new Runnable {
2         override def run(): Unit = Utils.tryLogNonFatalError {
3             self.send(CheckForWorkerTimeOut)
4         }
5     }, 0, WORKER_TIMEOUT_MS, TimeUnit.MILLISECONDS)

```

也专门起了一个线程给自己发送CheckForWorkerTimeOut信号。Master receive到CheckForWorkerTimeOut信号后：

```

1  case CheckForWorkerTimeOut =>
2  timeoutDeadWorkers()

```

调用timeoutDeadWorkers:

```

1  private def timeoutDeadWorkers() {
2      val currentTime = System.currentTimeMillis()
3      // 过滤出 最后收到心跳的时间 < 现在的时间 - worker心跳间隔的worker
4      val toRemove = workers.filter(_.lastHeartbeat < currentTime - WORKER_TIMEOUT_MS).toArray
5      // 遍历这些worker
6      for (worker <- toRemove) {
7          // 若WorkerInfo 状态不为 DEAD
8          if (worker.state != WorkerState.DEAD) {
9              logWarning("Removing %s because we got no heartbeat in %d seconds".format(
10                  worker.id, WORKER_TIMEOUT_MS / 1000))
11                  // 调用removeWorker
12                  removeWorker(worker)
13              }
14              // 若WorkerInfo 状态为 DEAD
15              else {
16                  // 等待足够长的时间后,
17                  // 再将它从workers列表中移除 :
18                  // 最后收到心跳的时间 < 现在的时间 - worker心跳间隔 × REAPER_ITERATIONS
19                  // REAPER_ITERATIONS 由 spark.dead.worker.persistence 参数设置,
20                  // 默认为 15
21                  if (worker.lastHeartbeat < currentTime - ((REAPER_ITERATIONS + 1) * WORKER_TIMEOUT_MS)) {
22                      workers -= worker
23                  }
24              }
25          }
26      }

```

Master.removeWorker

```

1  private def removeWorker(worker: WorkerInfo) {
2      logInfo("Removing worker " + worker.id + " on " + worker.host + ":" + worker.port)
3      // 标志worker状态为DEAD
4      worker.setState(WorkerState.DEAD)
5      // 移除各个缓存
6      idToWorker -= worker.id
7      addressToWorker -= worker.endpoint.address
8      if (reverseProxy) {
9          webUi.removeProxyTargets(worker.id)
10     }
11     for (exec <- worker.executors.values) {
12         logInfo("Telling app of lost executor: " + exec.id)
13         // 向使用该executor的app,
14         // 发送ExecutorUpdated信号
15         exec.application.driver.send(ExecutorUpdated(
16             exec.id, ExecutorState.LOST, Some("worker lost"), None, workerLost = true))
17         // 标志executor状态为LOST
18         exec.state = ExecutorState.LOST
19         // 将executor从app信息中移除
20         exec.application.removeExecutor(exec)
21     }
22     for (driver <- worker.drivers.values) {
23         // 重启 或移除 Driver
24         if (driver.desc.supervise) {
25             logInfo(s"Re-launching ${driver.id}")
26             relaunchDriver(driver)
27         } else {
28             logInfo(s"Not re-launching ${driver.id} because it was not supervised")
29             removeDriver(driver.id, DriverState.ERROR, None)

```

```

30     }
31 }
32 // 从持久化引擎中移除
33 persistenceEngine.removeWorker(worker)
34 }

```

Master.removeDriver

```

1  private def removeDriver(
2      driverId: String,
3      finalState: DriverState,
4      exception: Option[Exception]) {
5
6      drivers.find(d => d.id == driverId) match {
7          case Some(driver) =>
8              logInfo(s"Removing driver: $driverId")
9              // 从driver列表中移除
10             drivers -= driver
11             if (completedDrivers.size >= RETAINED_DRIVERS) {
12                 val toRemove = math.max(RETAINED_DRIVERS / 10, 1)
13                 completedDrivers.trimStart(toRemove)
14             }
15             // 加入到completedDrivers列表
16             completedDrivers += driver
17             // 从持久化引擎中移除
18             persistenceEngine.removeDriver(driver)
19             // 标志状态
20             driver.state = finalState
21             driver.exception = exception
22             // 将这个driver注册过的worker,
23             // 移除上面的driver
24             driver.worker.foreach(w => w.removeDriver(driver))
25             // 重新调度
26             schedule()
27         case None =>
28             logWarning(s"Asked to remove unknown driver: $driverId")
29     }
30 }
31 }

```

Master

接下来我们来讲讲Master的容错及HA。在之前的Master代码中出现了持久化引擎persistenceEngine的对象，其实它就是实现Master的容错及HA的关键。我们先来看看Master.osStart中，会根据RECOVERY_MODE，来生成持久化引擎persistenceEngine和选举代理leaderElectionAgent。

```

1  val (persistenceEngine_, leaderElectionAgent_) = RECOVERY_MODE match {
2      case "ZOOKEEPER" =>
3          logInfo("Persisting recovery state to ZooKeeper")
4          val zkFactory =
5              new ZooKeeperRecoveryModeFactory(conf, serializer)
6          (zkFactory.createPersistenceEngine(), zkFactory.createLeaderElectionAgent(this))
7      case "FILESYSTEM" =>
8          val fsFactory =
9              new FileSystemRecoveryModeFactory(conf, serializer)
10         (fsFactory.createPersistenceEngine(), fsFactory.createLeaderElectionAgent(this))
11      case "CUSTOM" =>
12          // 用户自定义机制
13          val clazz = Utils.classForName(conf.get("spark.deploy.recoveryMode.factory"))
14          val factory = clazz.getConstructor(classOf[SparkConf], classOf[Serializer])
15              .newInstance(conf, serializer)
16              .asInstanceOf[StandaloneRecoveryModeFactory]

```

```

17         (factory.createPersistenceEngine(), factory.createLeaderElectionAgent(this))
18     case _ =>
19         // 不做持久化
20         (new BlackHolePersistenceEngine(), new MonarchyLeaderAgent(this))
21     }
22     persistenceEngine = persistenceEngine_
23     leaderElectionAgent = leaderElectionAgent_

```

RECOVERY_MODE由spark.deploy.recoveryMode配置，默认为NONE：

```

1     private val RECOVERY_MODE = conf.get("spark.deploy.recoveryMode", "NONE")

```

接下来，我们来深入讲解下FILESYSTEM和ZOOKEEPER这两种recoveryMode。

FILESYSTEM

FILESYSTEM recoveryMode下，集群的元数据信息会保存在本地文件系统。而Master启动后则会立即成为Active的Master。

```

1     case "FILESYSTEM" =>
2         val fsFactory =
3             new FileSystemRecoveryModeFactory(conf, serializer)
4             (fsFactory.createPersistenceEngine(), fsFactory.createLeaderElectionAgent(this))

```

FileSystemRecoveryModeFactory会生成两个对象，一个是FileSystemPersistenceEngine，一个是MonarchyLeaderAgent：

```

1 private[master] class FileSystemRecoveryModeFactory(conf: SparkConf, serializer: Serializer)
2     extends StandaloneRecoveryModeFactory(conf, serializer) with Logging {
3
4     val RECOVERY_DIR = conf.get("spark.deploy.recoveryDirectory", "")
5
6     def createPersistenceEngine(): PersistenceEngine = {
7         logInfo("Persisting recovery state to directory: " + RECOVERY_DIR)
8         new FileSystemPersistenceEngine(RECOVERY_DIR, serializer)
9     }
10
11     def createLeaderElectionAgent(master: LeaderElectable): LeaderElectionAgent = {
12         new MonarchyLeaderAgent(master)
13     }
14 }

```

FileSystemRecoveryModeFactory

我们先来讲解下FileSystemRecoveryModeFactory：

```

1 private[master] class FileSystemPersistenceEngine(
2     val dir: String,
3     val serializer: Serializer)
4     extends PersistenceEngine with Logging {
5
6     // 新建一个目录
7     new File(dir).mkdir()
8
9     // 持久化对象，
10    // 将对象序列化的写入到文件
11    override def persist(name: String, obj: Object): Unit = {
12        serializeIntoFile(new File(dir + File.separator + name), obj)
13    }
14
15    // 去持久化
16    override def unpersist(name: String): Unit = {
17        val f = new File(dir + File.separator + name)
18        if (!f.delete()) {

```

```

19     logWarning(s"Error deleting ${f.getPath()}")
20   }
21 }
22
23 // 读取,
24 // 根据文件名反序列化出
25 override def read[T: ClassTag](prefix: String): Seq[T] = {
26   val files = new File(dir).listFiles().filter(_.getName.startsWith(prefix))
27   files.map(deserializeFromFile[T])
28 }
29
30 // 序列化到文件的实现
31 private def serializeIntoFile(file: File, value: AnyRef) {
32   // 生成新文件
33   val created = file.createNewFile()
34   if (!created) { throw new IllegalStateException("Could not create file: " + file) }
35   // 输出文件流
36   val fileOut = new FileOutputStream(file)
37   var out: SerializationStream = null
38   Utils.tryWithSafeFinally {
39     // 根据输出文件流 生成 输出序列化流
40     out = serializer.newInstance().serializeStream(fileOut)
41     // 将值通过输出序列化流写入文件
42     out.writeObject(value)
43   } {
44
45     // 关闭输出文件流
46     fileOut.close()
47     if (out != null) {
48       out.close()
49     }
50   }
51 }
52
53 // 从文件反序列化的实现
54 private def deserializeFromFile[T](file: File)(implicit m: ClassTag[T]): T = {
55   // 输入文件流
56   val fileIn = new FileInputStream(file)
57   var in: DeserializationStream = null
58   try {
59     // 根据输入文件流 生成 输入序列化流
60     in = serializer.newInstance().deserializeStream(fileIn)
61     // 从文件反序列化读取对象
62     in.readObject[T]()
63   } finally {
64     // 关闭输入文件流
65     fileIn.close()
66     if (in != null) {
67       in.close()
68     }
69   }
70 }
71
72 }

```

MonarchyLeaderAgent

```

1 @DeveloperApi
2 trait LeaderElectionAgent {
3   val masterInstance: LeaderElectable
4   def stop() {}
5 }
6
7 @DeveloperApi

```



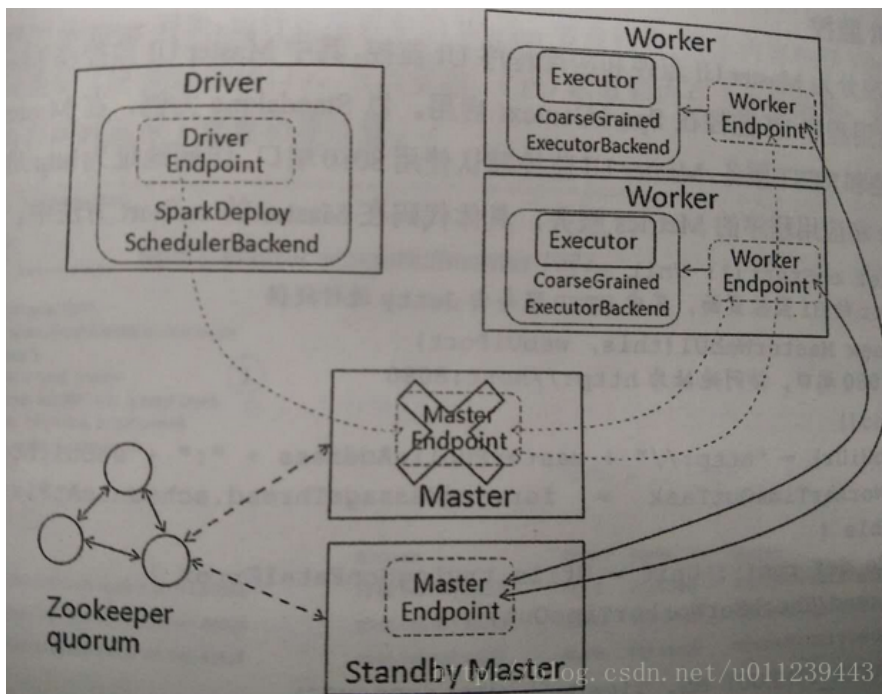
```

8  trait LeaderElectable {
9      def electedLeader(): Unit
10     def revokedLeadership(): Unit
11 }
12
13 // 选举代理的单点实现
14 // 总是启动最初的Master
15 private[spark] class MonarchyLeaderAgent(val masterInstance: LeaderElectable)
16     extends LeaderElectionAgent {
17     masterInstance.electedLeader()
18 }

```

ZOOKEEPER

ZOOKEEPER recoveryMode下，集群的元数据信息会保存在ZooKeeper中。ZooKeeper会在备份的Master中选举出新的Master，新的Master在启动后会从ZooKeeper中获取数据信息并且恢复这些数据。



```

1  case "ZOOKEEPER" =>
2      logInfo("Persisting recovery state to ZooKeeper")
3      val zkFactory =
4          new ZooKeeperRecoveryModeFactory(conf, serializer)
5      (zkFactory.createPersistenceEngine(), zkFactory.createLeaderElectionAgent(this))

```

ZooKeeperRecoveryModeFactory会生成两个对象，一个是ZooKeeperPersistenceEngine，一个是ZooKeeperLeaderElectionAgent：

```

1  private[master] class ZooKeeperRecoveryModeFactory(conf: SparkConf, serializer: Serializer)
2      extends StandaloneRecoveryModeFactory(conf, serializer) {
3
4      def createPersistenceEngine(): PersistenceEngine = {
5          new ZooKeeperPersistenceEngine(conf, serializer)
6      }
7
8      def createLeaderElectionAgent(master: LeaderElectable): LeaderElectionAgent = {
9          new ZooKeeperLeaderElectionAgent(master, conf)
10     }
11 }

```

ZooKeeperPersistenceEngine

我们先来讲解下ZooKeeperPersistenceEngine:

```
1 private[master] class ZooKeeperPersistenceEngine(conf: SparkConf, val serializer: Serializer)
2   extends PersistenceEngine
3   with Logging {
4
5   // 创建zk 及工作路径
6   private val WORKING_DIR = conf.get("spark.deploy.zookeeper.dir", "/spark") + "/master_status"
7   private val zk: CuratorFramework = SparkCuratorUtil.newClient(conf)
8
9   SparkCuratorUtil.mkdir(zk, WORKING_DIR)
10
11   // 持久化对象,
12   // 将对象序列化的写入到zk
13   override def persist(name: String, obj: Object): Unit = {
14     serializeIntoFile(WORKING_DIR + "/" + name, obj)
15   }
16
17   // 去持久化
18   override def unpersist(name: String): Unit = {
19     zk.delete().forPath(WORKING_DIR + "/" + name)
20   }
21
22   // 读取,
23   // 根据文件名反序列化出
24   override def read[T: ClassTag](prefix: String): Seq[T] = {
25     zk.getChildren.forPath(WORKING_DIR).asScala
26       .filter(_.startsWith(prefix)).flatMap(deserializeFromFile[T])
27   }
28
29   // 关闭zk
30   override def close() {
31     zk.close()
32   }
33
34   // 序列化到zk的实现
35   private def serializeIntoFile(path: String, value: AnyRef) {
36     // 序列化字节
37     val serialized = serializer.newInstance().serialize(value)
38     val bytes = new Array[Byte](serialized.remaining())
39     serialized.get(bytes)
40
41     // 写入到zk
42     zk.create().withMode(CreateMode.PERSISTENT).forPath(path, bytes)
43   }
44
45   // 从zk反序列化的实现
46   private def deserializeFromFile[T](filename: String)(implicit m: ClassTag[T]): Option[T] = {
47     // 从zk中得到数据
48     val fileData = zk.getData().forPath(WORKING_DIR + "/" + filename)
49     try {
50       // 反序列化
51       Some(serializer.newInstance().deserialize[T](ByteBuffer.wrap(fileData)))
52     } catch {
53       case e: Exception =>
54         logWarning("Exception while reading persisted file, deleting", e)
55         zk.delete().forPath(WORKING_DIR + "/" + filename)
56         None
57     }
58   }
59 }
```

ZooKeeperLeaderElectionAgent

ZooKeeperLeaderElectionAgent被创建会调用start:

```
1 private def start() {  
2     logInfo("Starting ZooKeeper LeaderElection agent")  
3     zk = SparkCuratorUtil.newClient(conf)  
4     leaderLatch = new LeaderLatch(zk, WORKING_DIR)  
5     leaderLatch.addListener(this)  
6     leaderLatch.start()  
7 }
```

leaderLatch.start(), 启动了leader的竞争与选举。涉及到的ZooKeeper选举实现, 已不在Spark源码范畴, 所以在这不再讲解。

总结

- **Executor**退出: 向**worker**发送**ExecutorStateChanged**信号; **worker**接收到信号后向**Master**发送**executorStateChanged**信号并释放该**Executor**资源; **Master**收到信号后, 改变该**Executor**状态, 移除**Web UI**上该**Executor**的信息, 若重试次数达到最大次数, 则移除该**Application**, 否则重新调度。
- **Worker**退出: **ShutdownHookManager**会调用**killProcess**杀死该所有的**executor**; **Master**利用心跳超时机制, 得知**Worker**退出, 改变该**Worker**状态, 将该**Worker**上的**Executor**从**Application**信息中移除, 将该**Worker**上的**driver**重启或移除, 从持久化引擎中移除该**Worker**。
- **Master**退出: **FILESYSTEM recoveryMode**下, 集群的元数据信息会保存在本地文件系统, 而**Master**启动后则会立即成为Active的**Master**; **ZOOKEEPER recoveryMode**下, 集群的元数据信息会保存在**ZooKeeper**中, **ZooKeeper**会在备份的**Master**中选举出新的**Master**, 新的**Master**在启动后会从**ZooKeeper**中获取数据信息并且恢复这些数据; 除此之外还有用户自定义的恢复机制和不做持久化的机制。